

# etl\_pipeline

August 5, 2025

Setup and imports

```
[4]: import logging
import requests
from datetime import datetime
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, udf, lit, current_timestamp
from pyspark.sql.types import DoubleType
from pyspark.sql.functions import to_date
import shutil
import os
```

Spark initialization

```
[5]: spark = SparkSession.builder \
    .appName("AdvancedETL") \
    .config("spark.jars", "/app/jars/mssql-jdbc-12.10.1.jre11.jar") \
    .getOrCreate()
```

Logs handling

```
[6]: logging.basicConfig(
    level=logging.INFO,
    format="%(asctime)s [%(levelname)s] %(message)s",
    handlers=[
        logging.FileHandler("/home/jovyan/etl_pipeline_error_log.log"),
        logging.StreamHandler()
    ]
)
```

Load source data

```
[7]: sales_df = spark.read.option("header", True).csv("data/sales_data_2.csv")
sales_df.show(5)

product_df = spark.read.option("header", True).csv("data/product_reference_2.
↪csv")
product_df.show(5)
```

OrderID	ProductID	SaleAmount	OrderDate	Region	CustomerID	Discount	Currency
1001	P50	299.99	01/05/2023	East	C100	0.1	USD
1002	P72	NULL	01/05/2023	West	C101	NULL	EUR
1003	P50	-10.0	01-06-2023	East	C100	0.05	GBP
1001	P50	299.99	01/05/2023	East	C100	0.1	USD
1004	P99	150.0	NULL	South	C102	0.2	USD

only showing top 5 rows

ProductID	ProductName	Category
P50	Wireless Mouse	Electronics
P72	Laptop Backpack	Accessories
P99	USB Hub	Electronics
P12	Notebook Stationery	Office Supplies
P88	Monitor Stand	Office Supplies

only showing top 5 rows

Null handling

```
[8]: sales_df.filter(col("SaleAmount").isNull()).show() # Check rows where
      ↳SaleAmount is null

sales_df.filter(col("OrderDate").isNull()).show() # Check rows where
      ↳OrderDate is null

sales_df.filter((col("SaleAmount").isNotNull()) & (col("OrderDate").
      ↳isNotNull())).show(5)
```

OrderID	ProductID	SaleAmount	OrderDate	Region	CustomerID	Discount	Currency
1002	P72	NULL	01/05/2023	West	C101	NULL	EUR

OrderID	ProductID	SaleAmount	OrderDate	Region	CustomerID	Discount	Currency
1004	P99	150.0	NULL	South	C102	0.2	USD

OrderID	ProductID	SaleAmount	OrderDate	Region	CustomerID	Discount	Currency
---------	-----------	------------	-----------	--------	------------	----------	----------

1001	P50	299.99	01/05/2023	East	C100	0.1	USD
1003	P50	-10.0	01-06-2023	East	C100	0.05	GBP
1001	P50	299.99	01/05/2023	East	C100	0.1	USD
1005	PX1	89.5	01/07/2023	North	NULL	0.0	USD
1006	P72	200.0	2023-13-01	West	C101	0.15	EUR

only showing top 5 rows

Duplicate removal

```
[9]: sales_df_clean = (
    sales_df
    .dropna(subset=["SaleAmount", "OrderDate"])
    .dropDuplicates(["OrderID"])
    .withColumn("OrderDateParsed", to_date("OrderDate", "MM/dd/yyyy"))
    .filter(
        (col("SaleAmount").cast("double").isNotNull()) &
        (col("OrderDateParsed").isNotNull())
    )
)

sales_df_clean.show(5)
```

OrderID	ProductID	SaleAmount	OrderDate	Region	CustomerID	Discount	Currency	OrderDateParsed
1001	P50	299.99	01/05/2023	East	C100	0.1	USD	
1005	PX1	89.5	01/07/2023	North	NULL	0.0	USD	
1007	P12	120.0	01/05/2023	East	C105	0.1	GBP	
1008	P88	300.0	02/05/2023	North	C106	0.0	USD	
1009	P77	0.0	03/05/2023	South	C107	NULL	USD	

only showing top 5 rows

Lookup: Join with product reference

```
[10]: enriched_df = sales_df_clean.join(product_df, on="ProductID", how="left")
print(f"[INFO] enriched_df row count: {enriched_df.count()}")
enriched_df.show(5)
```

```
[INFO] enriched_df row count: 15
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|ProductID|OrderID|SaleAmount|
OrderDate|Region|CustomerID|Discount|Currency|OrderDateParsed|
ProductName|      Category|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|      P50|  1001|    299.99|01/05/2023|  East|      C100|    0.1|    USD|
2023-01-05|    Wireless Mouse|    Electronics|
|      PX1|  1005|     89.5|01/07/2023| North|     NULL|    0.0|    USD|
2023-01-07|          NULL|          NULL|
|      P12|  1007|    120.0|01/05/2023|  East|      C105|    0.1|    GBP|
2023-01-05|Notebook Stationery|Office Supplies|
|      P88|  1008|    300.0|02/05/2023| North|      C106|    0.0|    USD|
2023-02-05|    Monitor Stand|Office Supplies|
|      P77|  1009|     0.0|03/05/2023| South|      C107|    NULL|    USD|
2023-03-05| Portable Speaker|    Electronics|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Currency conversion via API

```
[11]: def get_exchange_rates():
    try:
        url = "https://api.exchangerate-api.com/v4/latest/USD"
        response = requests.get(url)
        return response.json().get("rates", {})
    except Exception as e:
        logging.error(f"Exchange rate API failed: {e}")
        return {"EUR": 1.0, "GBP": 1.0}

exchange_rates = get_exchange_rates()
broadcast_rates = spark.sparkContext.broadcast(exchange_rates)

@udf(DoubleType())
def convert_to_usd(amount, currency):
    try:
        rate = broadcast_rates.value.get(currency, 1.0)
        return float(amount) / float(rate)
    except Exception as e:
```

```

        logging.error(f"Conversion error: amount={amount}, currency={currency},\n
↳error={e}")
        return None

converted_df = enriched_df.withColumn("SaleAmountUSD",\n
↳convert_to_usd(col("SaleAmount"), col("Currency")))
print(f"[INFO] converted_df row count: {converted_df.count()}")
converted_df.show(5)

```

[INFO] converted\_df row count: 15

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
|ProductID|OrderID|SaleAmount|
OrderDate|Region|CustomerID|Discount|Currency|OrderDateParsed|
ProductName|      Category|      SaleAmountUSD|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
|      P50|   1001|      299.99|01/05/2023|  East|      C100|      0.1|      USD|
2023-01-05|      Wireless Mouse|      Electronics|      299.99|
|      PX1|   1005|      89.5|01/07/2023| North|      NULL|      0.0|      USD|
2023-01-07|      NULL|      NULL|      89.5|
|      P12|   1007|      120.0|01/05/2023|  East|      C105|      0.1|      GBP|
2023-01-05|Notebook Stationery|Office Supplies|159.5744680851064|
|      P88|   1008|      300.0|02/05/2023| North|      C106|      0.0|      USD|
2023-02-05|      Monitor Stand|Office Supplies|      300.0|
|      P77|   1009|      0.0|03/05/2023| South|      C107|      NULL|      USD|
2023-03-05| Portable Speaker|      Electronics|      0.0|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 5 rows

Logging conversion info

```

[12]: conversion_log_df = converted_df.withColumn("ConversionTime",\n
↳current_timestamp()) \
      .select("OrderID", "Currency", "SaleAmount", "SaleAmountUSD",\n
↳"ConversionTime")

log_path = "/app/logs/conversion_log"

# Clean entire log directory if it exists
if os.path.exists(log_path):
    try:
        shutil.rmtree(log_path) # deletes folder and contents
        print(f"Deleted old log directory at {log_path}")
    except Exception as e:

```

```

        print(f"[WARN] Failed to delete log directory: {e}")

# Spark will create this folder fresh
conversion_log_df.coalesce(1).write \
    .mode("overwrite") \
    .option("header", True) \
    .csv(log_path)

```

Deleted old log directory at /app/logs/conversion\_log

Error handling with trashold

```

[13]: error_df = converted_df.filter(col("SaleAmountUSD").isNull())
error_df = error_df.withColumn("ErrorReason", lit("Invalid currency or_
    ↳amount")) \
        .withColumn("RejectedAt", current_timestamp())
error_df.write.mode("overwrite").option("header", True).csv("rejected/
    ↳rejected_records.csv")
error_df.show(5)

error_rate = error_df.count() / converted_df.count()
if error_rate > 0.05:
    raise Exception(f"[ERROR] Rejected records exceed 5% threshold_
    ↳({error_rate*100:.2f}%)")

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
|ProductID|OrderID|SaleAmount|OrderDate|Region|CustomerID|Discount|Currency|OrderDateParsed|Product
Name|Category|SaleAmountUSD|ErrorReason|RejectedAt|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+

```

Final clean data

```

[14]: final_df = converted_df.filter(col("SaleAmountUSD").isNotNull())

```

Write to SQL Database

```

[15]: jdbc_url = "jdbc:sqlserver://host.docker.internal:1433;databaseName=SalesDB;
    ↳encrypt=true;trustServerCertificate=true"

db_props = {
    "user": "sa",
    "password": "qwe123!@#",
    "driver": "com.microsoft.sqlserver.jdbc.SQLServerDriver"
}

```

```
}  
  
final_df.write.jdbc(url=jdbc_url, table="SalesEnriched", mode="append",  
    ↪properties=db_props)
```

Write rejected records to SQL for tracking

```
[16]: error_df.write.jdbc(url=jdbc_url, table="RejectedRecords", mode="append",  
    ↪properties=db_props)
```